



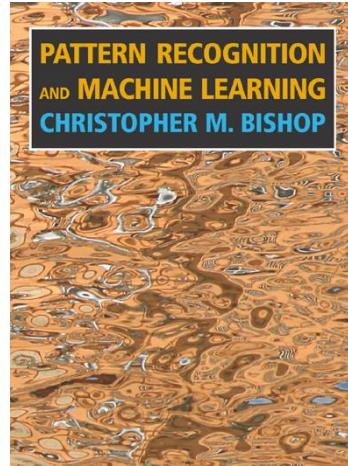
# Linear Regression in Pharmacology

Mohammad Ghazi Vakili, PhD  
Pharmacology Department

*Applying Statistical Learning to Drug Discovery and Clinical Data*



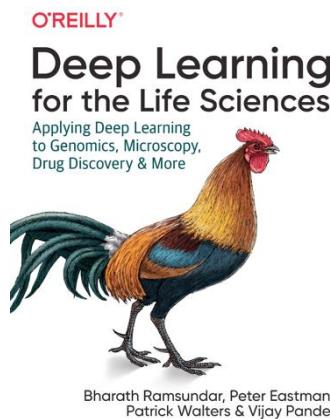
# Machine learning algorithms



## Pattern Recognition and Machine Learning

by Christopher M. Bishop

✓ Chapter 3: Linear Model for Regression



## Deep Learning for the Life Sciences

by Bharath Ramsundar, Peter Eastman, and Vijay S. Pande

✓ Chapter 2: Introduction to Deep Learning

Prerequisite: Calculus I



# Course structure

- **Weekly Lectures:**  
Introduce core ML concepts in Pharmacology applications
- **Python Lab:**  
Weekly Hands-on coding to implement algorithms and visualize data
- **Quizzes (30%):**  
Focus on understanding *when and why* to use each ML algorithm
- **Project (40%):** :  
Build an end-to-end ML pipeline for a real Pharmacology dataset
- **Exam (30%):** :  
Assess both conceptual understanding and practical application



# Machine learning algorithms

## What we will cover today's lecture

- Why Machine Learning
- QSAR
- QSAR and Linear Regression
- Example and Exercise



# Why LR in Pharmacology

- Pharmacology is data-rich
- Need to predict potency, toxicity, and Pharmacokinetics.
- Machine Learning (ML) enables:
  - Discovery of hidden patterns
  - Quantitative prediction of biological activity
  - Linking chemical structure → pharmacological behavior
- Next is Deep Learning



# What is QSAR

**QSAR (Quantitative Structure–Activity Relationship):**  
A mathematical model that relates molecular **structure** to biological **activity**

- Hansch & Fujita (1964):
- Led to famous **multivariate regression models** in pharmacology
- First linear equation:

$$\log\left(\frac{1}{c}\right) = a * \log P + b * \sigma + c * MW + d$$



# What is QSAR

## QSAR in Practice: A Simple Example

### Predicting Inhibitor Potency ( $\text{pIC}_{50}$ ) for a Protease Target

Compound	Mol. Wt. (Da)	LogP	$\text{pIC}_{50}$ (activity)
A	250	2.1	6.2
B	280	2.8	6.9
C	310	3.2	7.1
D	340	3.9	7.5
E	320	3.2	??

Question: Can we build a model to predict  $\text{pIC}_{50}$  for new compounds?



# Machine learning algorithms

## Supervised learning

- Data include **input features** and known **output labels**
  - Model learns a mapping:  $f(x) \rightarrow y$
- Goal: **predict or estimate outcomes for new data**
- Examples in pharmacology:
  - Predicting IC<sub>50</sub> (regression)
  - Classifying compounds as toxic/non-toxic (classification)

Algorithms:  
Linear Regression, Logistic  
Regression, Decision Trees, Neural  
Networks

## Unsupervised learning

- Data have **no predefined labels**
  - Model finds hidden structure or patterns
- Goal: **group or reduce data complexity**
- Examples in pharmacology:
  - Clustering compounds by chemical similarity
  - Grouping patients by response profile
  - Dimensionality reduction of molecular descriptors (e.g., PCA)

Algorithms:  
k-Means, Hierarchical Clustering,  
PCA, Autoencoders



# Mathematics Recap

## Vector and Matrix

$$x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

## Matrix

$$xA = [x_1 \quad x_2] \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} = [x_1 a_{11} + x_2 a_{21} \quad x_1 a_{12} + x_2 a_{22}]$$

$$B = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \quad (B)^{-1} = \frac{1}{ad-bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

$$B^T = \begin{bmatrix} a & c \\ b & d \end{bmatrix}$$



# Linear Regression

Chapter 3



# Linear Regression:

Training set: Data used to train the model

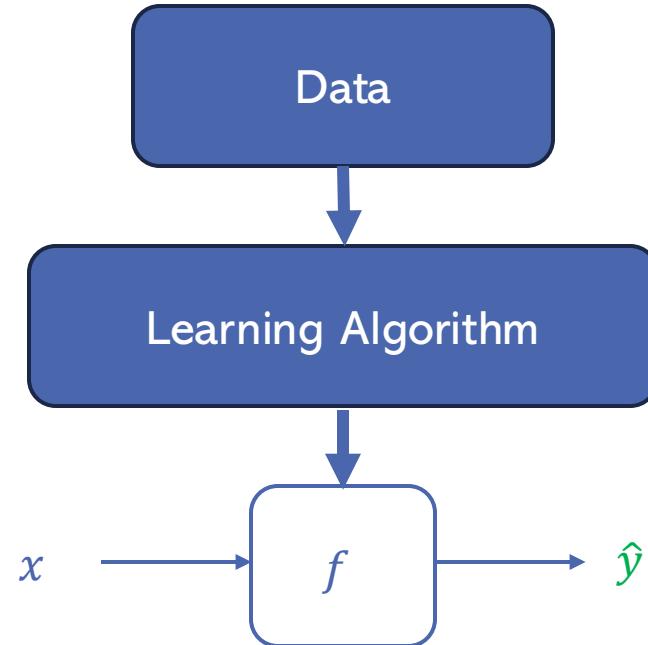
Compound	Mol. Wt. (Da)	LogP	pIC <sub>50</sub> (activity)
A	250	2.1	6.2
B	280	2.8	6.9
C	310	3.2	7.1
D	340	3.9	7.5
...	...	...	...
E	320	3.2	??

$x$  : Input (features)

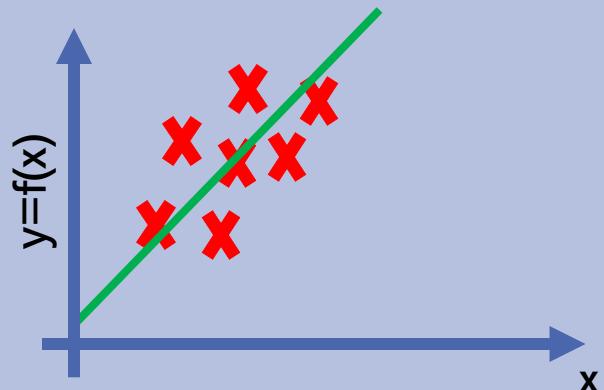
$y$ : output (target)



# Linear Regression



How to represent  $f$



One size variable

$$f_{w,b}(x) = wx + b$$

Multi variable

$$f_{w,b}(x) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 =$$

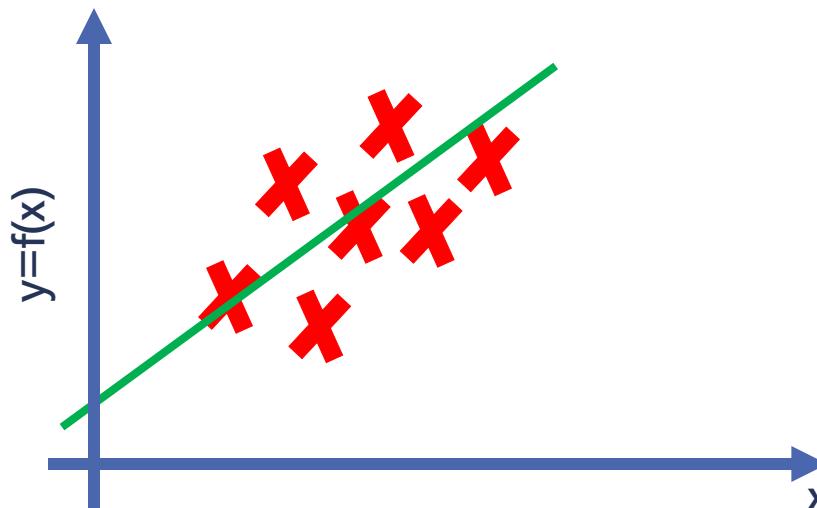
$$w_0 + \sum_{i=0}^3 w_i x_i$$



# Linear Regression: Model

Cost Function  $J(w, b)$

Measures the average squared difference between predicted ( $\hat{y}$ ) and actual ( $y$ ) values. The goal is to find parameters  $w, b$  that **minimize this error**.



$$\hat{y} = f_{w,b}(x^{(i)})$$

$$f_{w,b}(x) = w_0 + w x^{(i)}$$

$$J(w) = \frac{1}{2m} \sum_{i=0}^m (\hat{y}^{(i)} - y^{(i)})^2$$

With :

$$\hat{y} = f_{w,b}(x^{(i)}) = w_0 + w x^{(i)}$$

Find the parameters  $w$  and  $b$  that **minimize the cost function  $J(w, b)$**



# Closed Form (Normal Equation)

To minimize the  $J(w)$  take the derivative with respect to  $w$  and set it to zero:

$$\frac{\partial J}{\partial w} = \left(\frac{1}{m}\right) \Phi^T (\Phi w - y) = 0$$

Where:

$$\begin{aligned}\hat{y} &= wx + b \\ \hat{y} &= w_1x + w_0 = [w_0 \ w_1] \begin{bmatrix} 1 \\ x \end{bmatrix} = w\phi\end{aligned}$$

Simplify

$$\phi^T \phi w = \phi^T y$$

Now solve  $w$

$$w = (\phi^T \phi)^{-1} \phi^T y$$

This is called the **Normal Equation**, it gives the exact, closed-form solution.

Note:

- Requires invertible  $\phi^T \phi w$  (issue with multicollinearity).
- Gradient Decent will be used for higher dimension (Training)



# The Coefficient of Determination

Measures how well our regression line fits the data

$$R^2 = 1 - \left( \frac{SS_{res}}{SS_{tot}} \right)$$

$SS_{res} = \sum(y_i - \hat{y}_i)^2 \rightarrow$  unexplained variation (residuals)

$SS_{tot} = \sum(y_i - \bar{y}_i)^2 \rightarrow$  total variation in data

## Interpretation:

- $R^2 = 1 \rightarrow$  Perfect fit (model explains all variation)
- $R^2 = 0 \rightarrow$  Model explains none of the variation
- Example:  $R^2 = 0.85$  means the model explains **85% of the change** in the response.



# Recap

- QSAR

$$\hat{Y} = \log\left(\frac{1}{c}\right) = a * \log P + b * \sigma + c * MW + d$$

- Linear Regression

$$\hat{y} = w_1 x + w_0 = [w_0 \ w_1] \begin{bmatrix} 1 \\ x \end{bmatrix} = w\phi$$

$$J(w) = \frac{1}{2m} \sum_{i=0}^m (\hat{y}^{(i)} - y^{(i)})^2$$

- How to compute the coefficients

- Normal Equation

$$w = (\phi^T \phi)^{-1} \phi^T y \quad \text{and} \quad R^2 = 1 - \left( \frac{SS_{res}}{SS_{tot}} \right)$$



# Normal Function Example:

Fit Linear model to

No.	x	y
1	1	2
2	2	3
3	3	5

$$X = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{bmatrix} \quad y = \begin{bmatrix} 2 \\ 3 \\ 5 \end{bmatrix}$$

$$w = (X^T X)^{-1} X^T y$$

$$X^T X = \begin{bmatrix} 3 & 6 \\ 6 & 14 \end{bmatrix}$$

$$X^T y = \begin{bmatrix} 10 \\ 23 \end{bmatrix}$$

$$(X^T X)^{-1} = \frac{1}{6} \begin{bmatrix} 14 & -6 \\ -6 & 3 \end{bmatrix}$$

$$w = \begin{bmatrix} 0.33 \\ 1.5 \end{bmatrix}$$

$$y = 0.33 + 1.5x$$

$$R^2 = 0.964$$



# Lab Exercise 1

## Predicting Drug-Induced Heart Rate Response Using Multi-Feature Regression

	Dose (mg/kg)	Weight (kg)	Age (years)	Response ( $\Delta$ bpm)
1	3.62	74.3	42.49	9.05
2	7.66	56.82	37.21	21.94
3	6.12	52.6	62.29	17.46
4	5.19	87.96	41.05	16.55
5	2.09	88.63	37.64	2.71
6	2.09	82.34	49.42	5.51
7	1.41	62.18	31.34	-3.75
8	7.06	53.91	61.1	21.69
9	5.21	77.37	28.35	14.66

Develop a multiple linear regression model that predicts the heart rate response ( $\Delta$  bpm) based on the following independent variables:

- Dose (mg/kg): Drug dosage per kilogram of body weight
- Weight (kg): Subject's total body mass
- Age (years): Subject's age

The goal is to **quantify the contribution of each factor** to the observed pharmacological response and interpret the coefficients in a biological context.



# Lab Exercise 1

## Predicting Drug-Induced Heart Rate Response Using Multi-Feature Regression

	Dose (mg/kg)	Weight (kg)	Age (years)	Response (Δ bpm)
1	3.62	74.3	42.49	9.05
2	7.66	56.82	37.21	21.94
3	6.12	52.6	62.29	17.46
4	5.19	87.96	41.05	16.55
5	2.09	88.63	37.64	2.71
6	2.09	82.34	49.42	5.51
7	1.41	62.18	31.34	-3.75
8	7.06	53.91	61.1	21.69
9	5.21	77.37	28.35	14.66

$$\hat{Y} = w_0 + w_1 * \text{Dose} + w_2 * \text{Weight} + w_3 * \text{Age}$$

$$w = (\phi^T \phi)^{-1} \phi^T y$$

Solution:

$$\hat{Y} = -3.231 + 3.410 * \text{Dose} + 0.003 * \text{Weight} - 0.023 * \text{Age} \text{ and } R^2 = 0.99$$

[Link to Google Colab](#)



# Next session

- Gradient Decent
- Python
- Python implementation of the LR
  - Normal Equation
- Non-Linear LR
  - Gaussian (Radial Basis) Regression
  - Logarithmic and Exponential Models
  - Generalized Linear Models (GLMs)
  - Sigmoidal / Logistic Regression (Hill Equation)



# Appendix A



# QSAR Parameters

$$\log\left(\frac{1}{C}\right) = a * \log P + b * \sigma + c * MW + d$$

Symbol	Description	Pharmacological Interpretation
C	Molar concentration required to produce a defined biological effect	Inversely related to potency (lower C → higher activity)
$\log(1/C)$	Activity term (higher = more active)	Used as dependent variable (Y)
$\log P$	Partition coefficient (lipophilicity between octanol and water)	Measures hydrophobicity; affects membrane permeability and receptor binding
$\sigma$ (sigma)	Hammett substituent constant	Represents electronic effects of substituents on the aromatic ring (electron-withdrawing or -donating)
MW	Molecular weight	Captures steric or size-related contribution to activity
a, b, c	Regression coefficients	Quantify how each property influences activity (sign and magnitude indicate direction and strength)
d	Intercept	Baseline biological activity independent of descriptors



# Linear Regression: Model

We want to find the best-fitting line that predicts  $\hat{y}$  from  $x$ :

$$\begin{aligned}\hat{y} &= wx + b \\ \hat{y} &= w_1x + w_0 = [w_0 \ w_1] \begin{bmatrix} 1 \\ x \end{bmatrix}\end{aligned}$$

or, in vector/matrix form (for many samples and features):

$$\hat{y} = w\phi$$

Where

- $\phi$  = design matrix (rows = samples, columns = features, plus 1 for intercept),
- $w$  = weights (parameters),
- $\hat{y}$  = predicted values.



# Gradient Descent (When D or N is large)

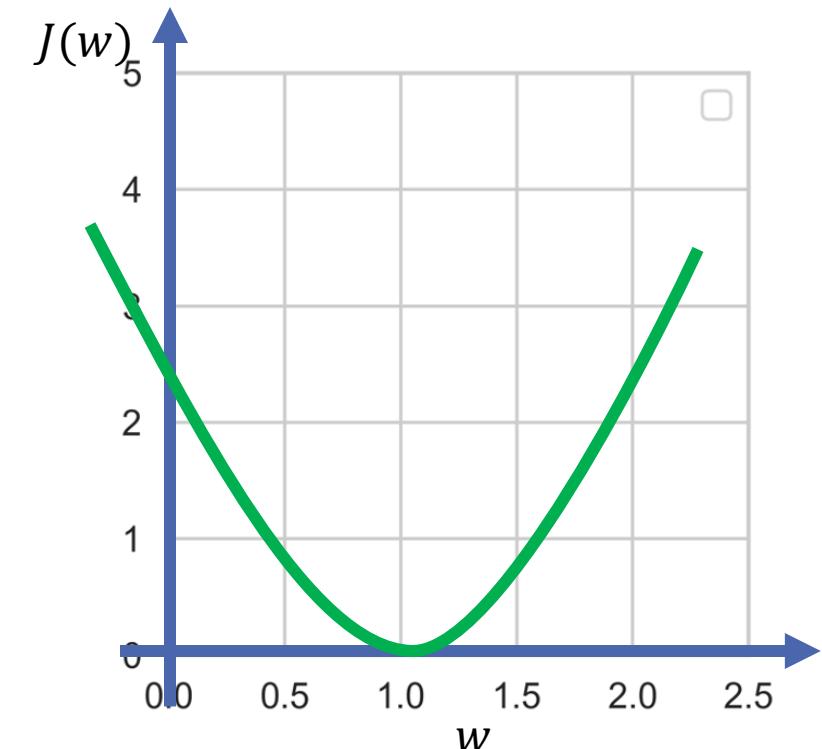
This **iterative optimization** steps is called Gradient Decent Instead of computing the  $(\phi^T \phi)^{-1} \phi^T y$ , it **updates weights step-by-step** by following the gradient (slope) of the cost function:

$$w := w - \alpha \frac{\partial J}{\partial w}$$

where  $\alpha$  is the **learning rate**.

$$\frac{\partial J}{\partial w} = \left( \frac{1}{m} \right) \sum (\hat{y}^{(i)} - y^{(i)}) x^{(i)}$$

- Scales to streaming data; no matrix inversion.
- Learning rate  $\alpha$  controls convergence speed/stability.
- Batch / stochastic / mini-batch variants.





# Gradient Descent (When D or N is large)

